

# ORCA

## Query optimization as a service

@addisonhuddy

**20 min**

- Little about query optimization**
- Introduce ORCA**
- ORCA Internals**
- Pairing: adding a transformation**
- ORCA Roadmap**

# Why care about query optimization?

- Turns queries (SQL, MapReduce, ...) into an execution plan
- Data growth > Processing growth
- So many optimizers!

# Why care about ORCA?

- **Modular**
- **Extensible**
- **Plugable**

=> **R&D test-bed**

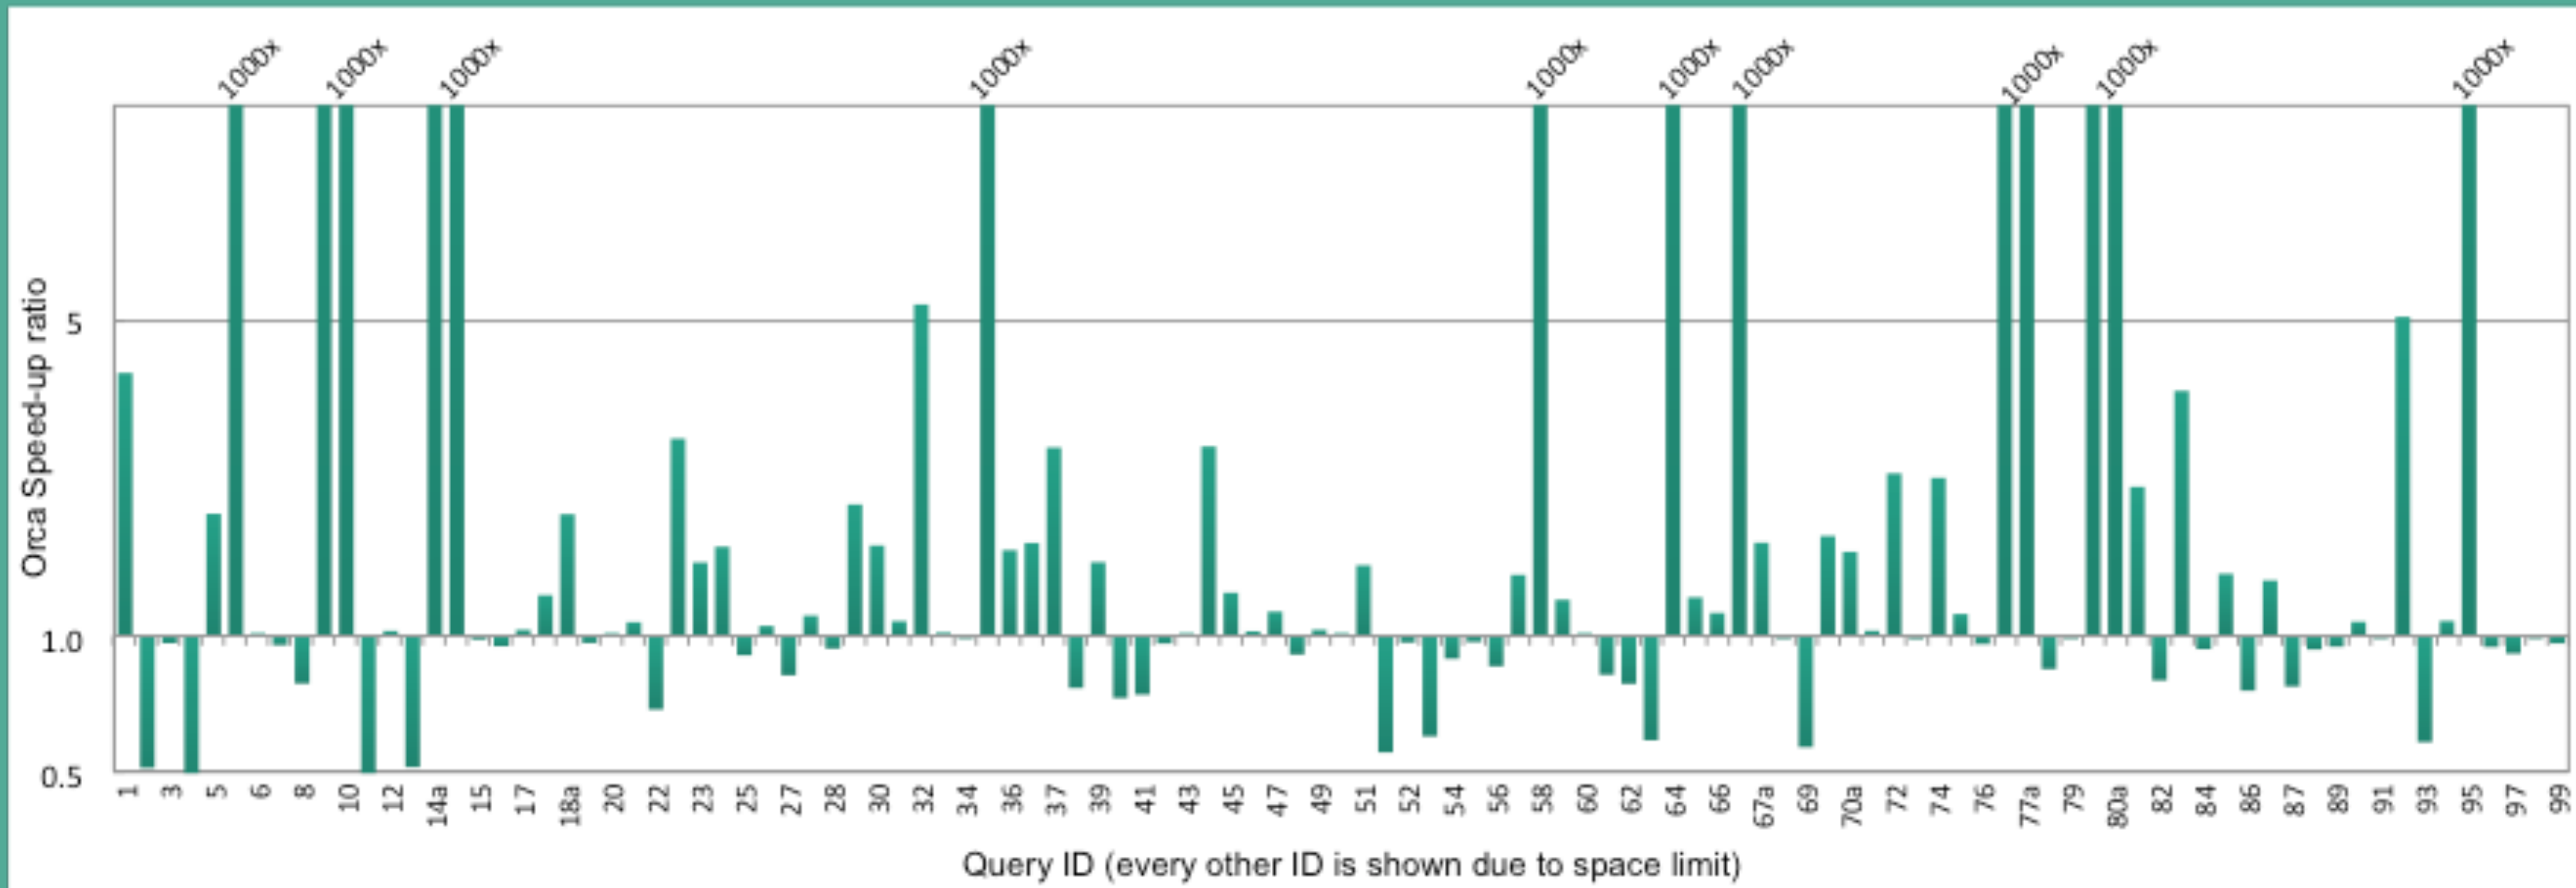


**HAWQ**



**GREENPLUM  
DATABASE**

# TPC-DS 5X Faster<sup>1</sup>



<sup>1</sup> © 2014 ACM, used with permission. Figure 3. TPC-DS performance testing results of Pivotal Greenplum with Pivotal Query Optimizer vs. Pivotal Greenplum with “planner” query optimizer.



# Greenplum Database

Massively-Parallel, Shared-Nothing Database based on PostgreSQL

© 2014 Greenplum Inc. <http://www.greenplum.com> [gpdb-dev@greenplum.com](mailto:gpdb-dev@greenplum.com)

# Now Open Source

 **Repositories**

 **People** 35

**Filters** ▾

 Find a repository...

## gpdb

Greenplum Database

Updated 40 minutes ago

PLpgSQL ★ 1,300 🍴 305

## gporca

A modular query optimizer for big data

Updated 22 hours ago


C++ ★ 22 🍴 11

## gpos

C++ ★ 2 🍴 1


## People

35 >





# What makes it so unique?





# Key Features

- **Smarter partition elimination**
- **Subquery unnesting**
- **Common table expressions (CTE)**
- **Join ordering**
- **Sort order optimization**
- **Skew awareness**

# Logical & || Physical Operators

**Pre-processing -> Exploration -> Implementation -> Optimization**

**ALL possible logical plans are turned into its physical operators**

**- Dr. Venkatesh "Venky" Raghavan**

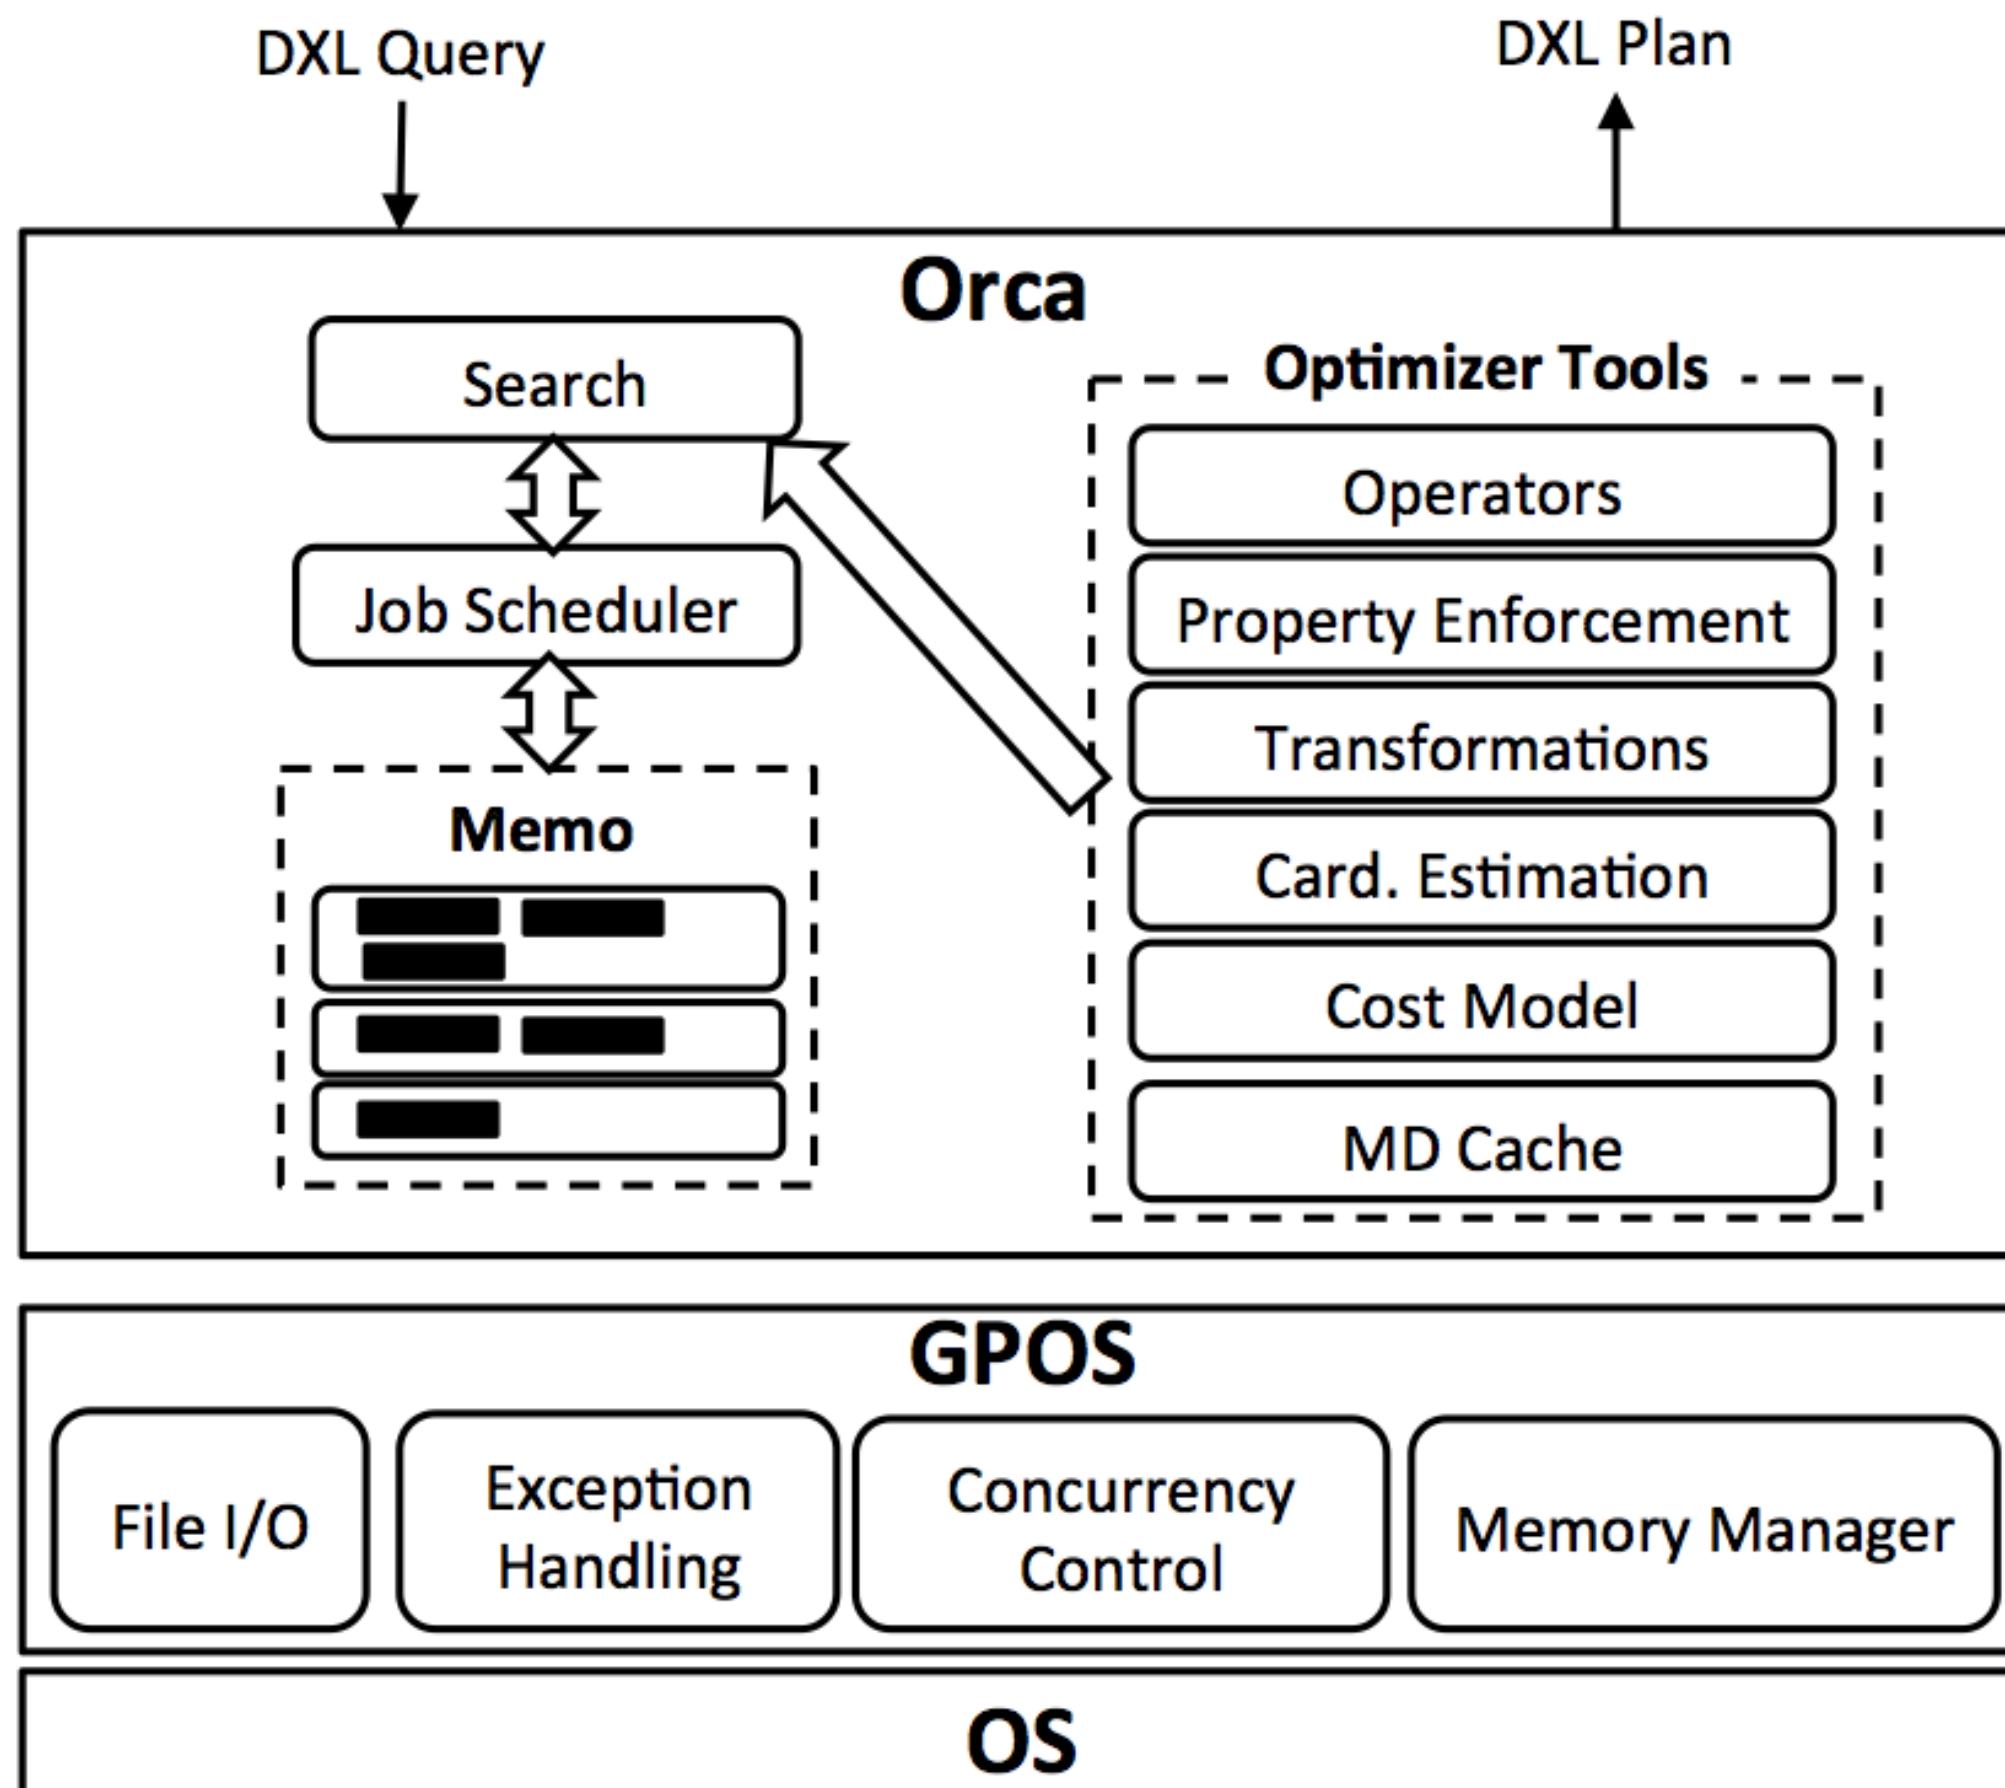
# Many Logical transformations

- **Join Ordering Algorithm**
- **Expand NAry Join Min Card Algorithm**
- **Expand NAry Join Dynamic Programming**
- **Select to Filter**
- **Select to IndexGet**
- **Simplify Select With Subquery**

**+77 more from what I could count on the plane**

1,000,000,000





# Let's Pair

# Idea

**Split an aggregate into a pair of local and global aggregate.**

```
SELECT sum(c) FROM foo GROUP BY b
```

```
CREATE TABLE foo (a int, b int, c int) distributed by (a);
```

# CXformSplitGbAgg

// HEADER FILES

~/orca/libgpopt/include/gpopt/xforms

// SOURCE FILES

~/orca/libgpopt/src/xforms

# Pattern

```
GPOS_NEW(pmp)
CExpression
(
    pmp,
    GPOS_NEW(pmp) CLogicalGbAgg(pmp),
    // logical aggregate operator
    GPOS_NEW(pmp) CExpression(pmp, GPOS_NEW(pmp) CPatternLeaf(pmp)),
    // relational child
    GPOS_NEW(pmp) CExpression(pmp, GPOS_NEW(pmp) CPatternTree(pmp))
    // scalar project list
));
```



# What?

```
...
-> Redistribute Motion 2:2 (slice1; segments: 2) (cost=0.00..431.00 rows=1 width=12)
    Hash Key: b
    Rows out: Avg 1.5 rows x 2 workers at destination. Max 2 rows (seg0) with 1.105 ms to end, start offset by 15 ms.
-> Result (cost=0.00..431.00 rows=1 width=12)
    Rows out: Avg 1.5 rows x 2 workers. Max 2 rows (seg1) with 0.273 ms to first row, 0.276 ms to end, start offset by 16 ms.
-> GroupAggregate (cost=0.00..431.00 rows=1 width=12)
    Group By: b
    Rows out: Avg 1.5 rows x 2 workers. Max 2 rows (seg1) with 0.272 ms to firstrow, 0.275 ms to end, start offset by 16 ms.
-> Sort (cost=0.00..431.00 rows=1 width=8)
    Sort Key: b
    Rows out: Avg 1.5 rows x 2 workers. Max 2 rows (seg1) with 0.260 ms to first row, 0.262 ms to end, start offset by 16 ms.
    Executor memory: 58K bytes avg, 58K bytes max (seg0).
    Work_mem used: 58K bytes avg, 58K bytes max (seg0). Workfile: (0 spilling, 0 reused)
-> Table Scan on foo2 (cost=0.00..431.00 rows=1 width=8)
    Rows out: Avg 1.5 rows x 2 workers. Max 2 rows (seg1) with 0.074 ms to first row, 0.075 ms to end, start offset by 16 ms.
```

# Pre-condition Check

```
// Compatibility function for splitting aggregates  
virtual  
BOOL FCompatible(CXform::EXformId exfid){  
    return (CXform::ExfSplitGbAgg != exfid);}
```

# The Actual Transformation

```
void Transform  
(  
    CXformContext *pxfctxt,  
    CXformResult *pxfres,  
    CExpression *pexpr  
) const;
```


# Register Transformation

```
void CXformFactory::Instantiate()  
{  
...  
Add(GPOS_NEW(m_pmp) CXformSplitGbAgg(m_pmp));  
...  
}
```

# What can't it do?



# Balance



# Improved performance for short running queries

# Not yet feature complete

- **External parameters**
- **Cubes**
- **Multiple grouping sets**
- **Inverse distribution functions**
- **Ordered aggregates**
- **Indexed expressions**

# PostgreSQL



# Four pieces of low hanging fruit

- 1. Distinguish between Physical and Logical**
- 2. Move expression evaluation inside ORCA**
- 3. ORCA assumes indexes are all forward access**
- 4. Constraint evaluation from a NOT NULL**



# Get Involved

**[github.com/greenplum-db](https://github.com/greenplum-db)**

**[gpdb-dev@greenplum.org](mailto:gpdb-dev@greenplum.org)**

**Pivotal Tracker: [bit.ly/1m1WGDn](https://bit.ly/1m1WGDn)**

**White Paper: [bit.ly/1ntrE8v](https://bit.ly/1ntrE8v)**

**[@addisonhuddy](#)**

# Slides

**Slides: [github.com/addisonhuddy/presentations](https://github.com/addisonhuddy/presentations)**